

Assignment#3 (ListScan, Histogram)

Due Date: Wednesday, Apr 29th 11:59PM

Steps for solving the assignment (refer to “Assignment Environment” lecture slides for details):

- Download the assignment from the Canvas and unzip it (now you have “A3” directory - you must have already done this to read this a3-instructions.pdf file).
- Build the source using CMake (in "build" directory that you create under "A3" directory).
- Copy datasets directly under “build” directory (under “A3/datasets” directory, we have two dataset directories, copy each of them to “A3/build”, e.g., copy contents of “A3/datasets/ListScan” to “A3/build/ListScan”).
- For each problem in the set:
 - Set the project properties as indicated in the below “Project Setup” instructions (e.g., set command line options).
 - Update the template code (template.cu) as indicated in the below “Coding” instructions to solve the problem.
 - Build your solution under Visual Studio (right click on the project and hit “Build”) and run it under Visual Studio (right click on the project and hit “Debug -> Start new instance”).
 - Make sure your executable (located under “Debug” directory that resides under “build” directory) works from the command line as indicated in the below “Command-line Execution” instructions.
 - Create a new directory by giving it the name of the problem (e.g., ListScan).
 - Copy the executable file (.exe file) and the updated template code (template.cu) to this newly created directory.
- Zip all newly created problem directories, name this ZIP file as YourLastName-YourFirstName-a#.zip (e.g., Doe-Jane-a3.zip).
- Build (using CMake and Visual Studio) and test your assignment in a **remotely accessible lab machine, ECS-GPU AVD** (see the syllabus for more info) and **create a TEXT (i.e., not a pdf, doc etc.) file called "readme-a3.txt"** that indicates whether your submission works on the remotely

accessible lab machine. You may also include additional information you want to share with the grader in this text file. **Make sure you have the executable files (.exe files) generated/tested on the remote lab machine in the zip file mentioned above.**

- Submit the **TWO files you have created above (zip file and txt file) SEPERATELY to Canvas** (i.e., do NOT place the txt file inside the zip file). You will receive the grader comments on your text file when grades are posted.

Each dataset directory provides **five test cases**. You should test your program with all the provided test cases. You can test individual test cases by updating the command line options provided in the below "Project Setup" instructions. For instance, to test test_case_1 of ListScan (located at "build/ListScan/Dataset\1") you should update the command line option provided below (which tests test_case_0 located at "build/ListScan/Dataset\0") by changing all subtexts that say "...Dataset\0\..." with "...Dataset\1\...".

If you want to quickly check whether all datasets of both problems included in A3 work from the command prompt, you can use the attached text file (A3_batch.txt) that lists the command-line execution lines for the datasets excluding the last part from the line that directs the output to result.txt. This allows us to display all results on the console instead of writing them to result.txt files. Copy/paste the contents of these text files to command prompt when you are under "build" directory to see whether your code returns correct results for all datasets of both problems.

Specific Problem Instructions:

Problem#1: ListScan

Objective

The objective of this problem is to implement a kernel to perform an **inclusive** parallel scan on a 1D list. The scan operator will be the addition (plus) operator. You should implement the **work efficient scan kernel** and handle input lists that have sizes larger than the maximum thread block size as discussed in the lecture notes. But you can assume that the input list length divided by twice the block size (which is set to 512) will not be more than twice the block size.

The boundary condition can be handled by filling "identity value (0 for sum)" into the shared memory of the last block when the length is not a multiple of the thread block size.

The device computation should be split into two kernels: one that computes scanned blocks and one that adds scanned block sums to all values of the scanned blocks.

Coding

Edit the template.cu to perform the following:

- allocate device memory and clear output on device by uncommenting the related line
- copy input host memory to device
- initialize thread block and kernel grid dimensions
- invoke CUDA kernels
- copy output device memory to host
- deallocate device memory
- write the CUDA kernels

Instructions about where to place each part of the code is demarcated by the `//@@` comment lines.

Project Setup

To test your program on `test_case_0`:

- Right click on the ListScan_Template project -> Properties -> Configuration Properties -> Debugging -> Command Arguments and enter the following:

```
-e .\ListScan\Dataset\0\output.raw  
-i .\ListScan\Dataset\0\input.raw  
-o .\ListScan\Dataset\0\myoutput.raw -t vector  
> .\ListScan\Dataset\0\result.txt
```

(all in one line - do not forget to put spaces between the sub-lines)

You will see the output of your execution in `result.txt` which resides under `build\ListScan\Dataset\0\`

Command-line Execution

The executable generated as a result of compiling the project (build\Debug\ListScan_Template.exe) can be run from the command-line using the following command (make sure you are in build directory):

```
.\Debug\ListScan_Template -e .\ListScan\Dataset\0\output.raw  
-i .\ListScan\Dataset\0\input.raw  
-o .\ListScan\Dataset\0\myoutput.raw -t vector  
> .\ListScan\Dataset\0\result.txt
```

(all in one line - do not forget to put spaces between the sub-lines)

Problem#2: Histogram

Objective

The purpose of this problem is to implement an efficient histogramming algorithm for an input array of integers within a given range. You should implement the **histogram kernel that uses privatization technique** as described in the class notes. Each integer will map into a single bin (in both shared and global memory), so the values will range from 0 to (NUM_BINS - 1). The histogram bins will use unsigned 32-bit counters that must be saturated at 127 (i.e., if the bin value is more than 127, make it equal to 127). The input length can be assumed to be at most 500000 elements. NUM_BINS is fixed at 4096.

The device computation should be split into two kernels: the first kernel should compute the histogram results without applying saturation and the second kernel should apply saturation to the histogram results that are produced by the first kernel. Note that these two stages could have been combined into a single kernel. However, in this assignment, we will keep them separated.

Coding

Edit the code in the code tab to perform the following:

- allocate device memory
- copy input host memory to device
- clear the bins on device using cudaMemset() (see the sample use of cudaMemset() in template.cu of ListScan)
- initialize thread block and kernel grid dimensions
- invoke CUDA kernels
- copy output device memory to host

- deallocate device memory
- write the CUDA kernels

Instructions about where to place each part of the code is demarcated by the `//@@` comment lines.

Project Setup

To test your program on `test_case_0`:

- Right click on the `Histogram_Template` project -> Properties -> Configuration Properties -> Debugging -> Command Arguments and enter the following:

```
-e .\Histogram\Dataset\0\output.raw
-i .\Histogram\Dataset\0\input.raw
-o .\Histogram\Dataset\0\myoutput.raw -t integral_vector
> .\Histogram\Dataset\0\result.txt
```

(all in one line - do not forget to put spaces between the sub-lines)

You will see the output of your execution in `result.txt` which resides under

`build\Histogram\Dataset\0\`

Command-line Execution

The executable generated as a result of compiling the project (`build\Debug\Histogram_Template.exe`) can be run from the command-line using the following command (make sure you are in `build` directory):

```
.\Debug\Histogram_Template
-e .\Histogram\Dataset\0\output.raw
-i .\Histogram\Dataset\0\input.raw
-o .\Histogram\Dataset\0\myoutput.raw -t integral_vector
> .\Histogram\Dataset\0\result.txt
```

(all in one line - do not forget to put spaces between the sub-lines)